

# Project Batcomputer

A working DevOps  
implementation for Machine  
Learning

**Ben Coleman**  
**Phil Harvey**

**@BenCodeGeek**  
**@CodeBeard**



*v0.0.3 (Beta)*

# Background

## Motivation

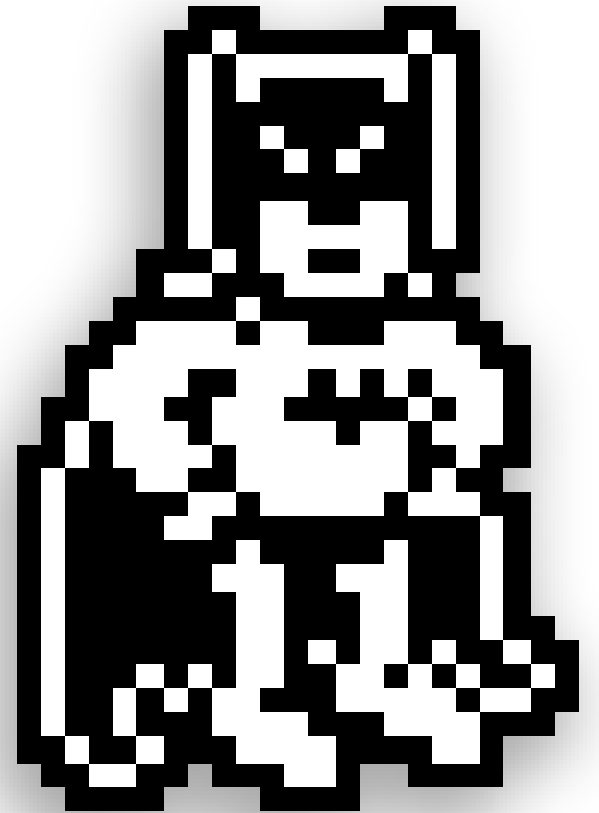
- Understand challenges in operationalisation of ML models
- Existing processes approaches deemed problematic
- “DevOps for AI”

## Why Batcomputer?

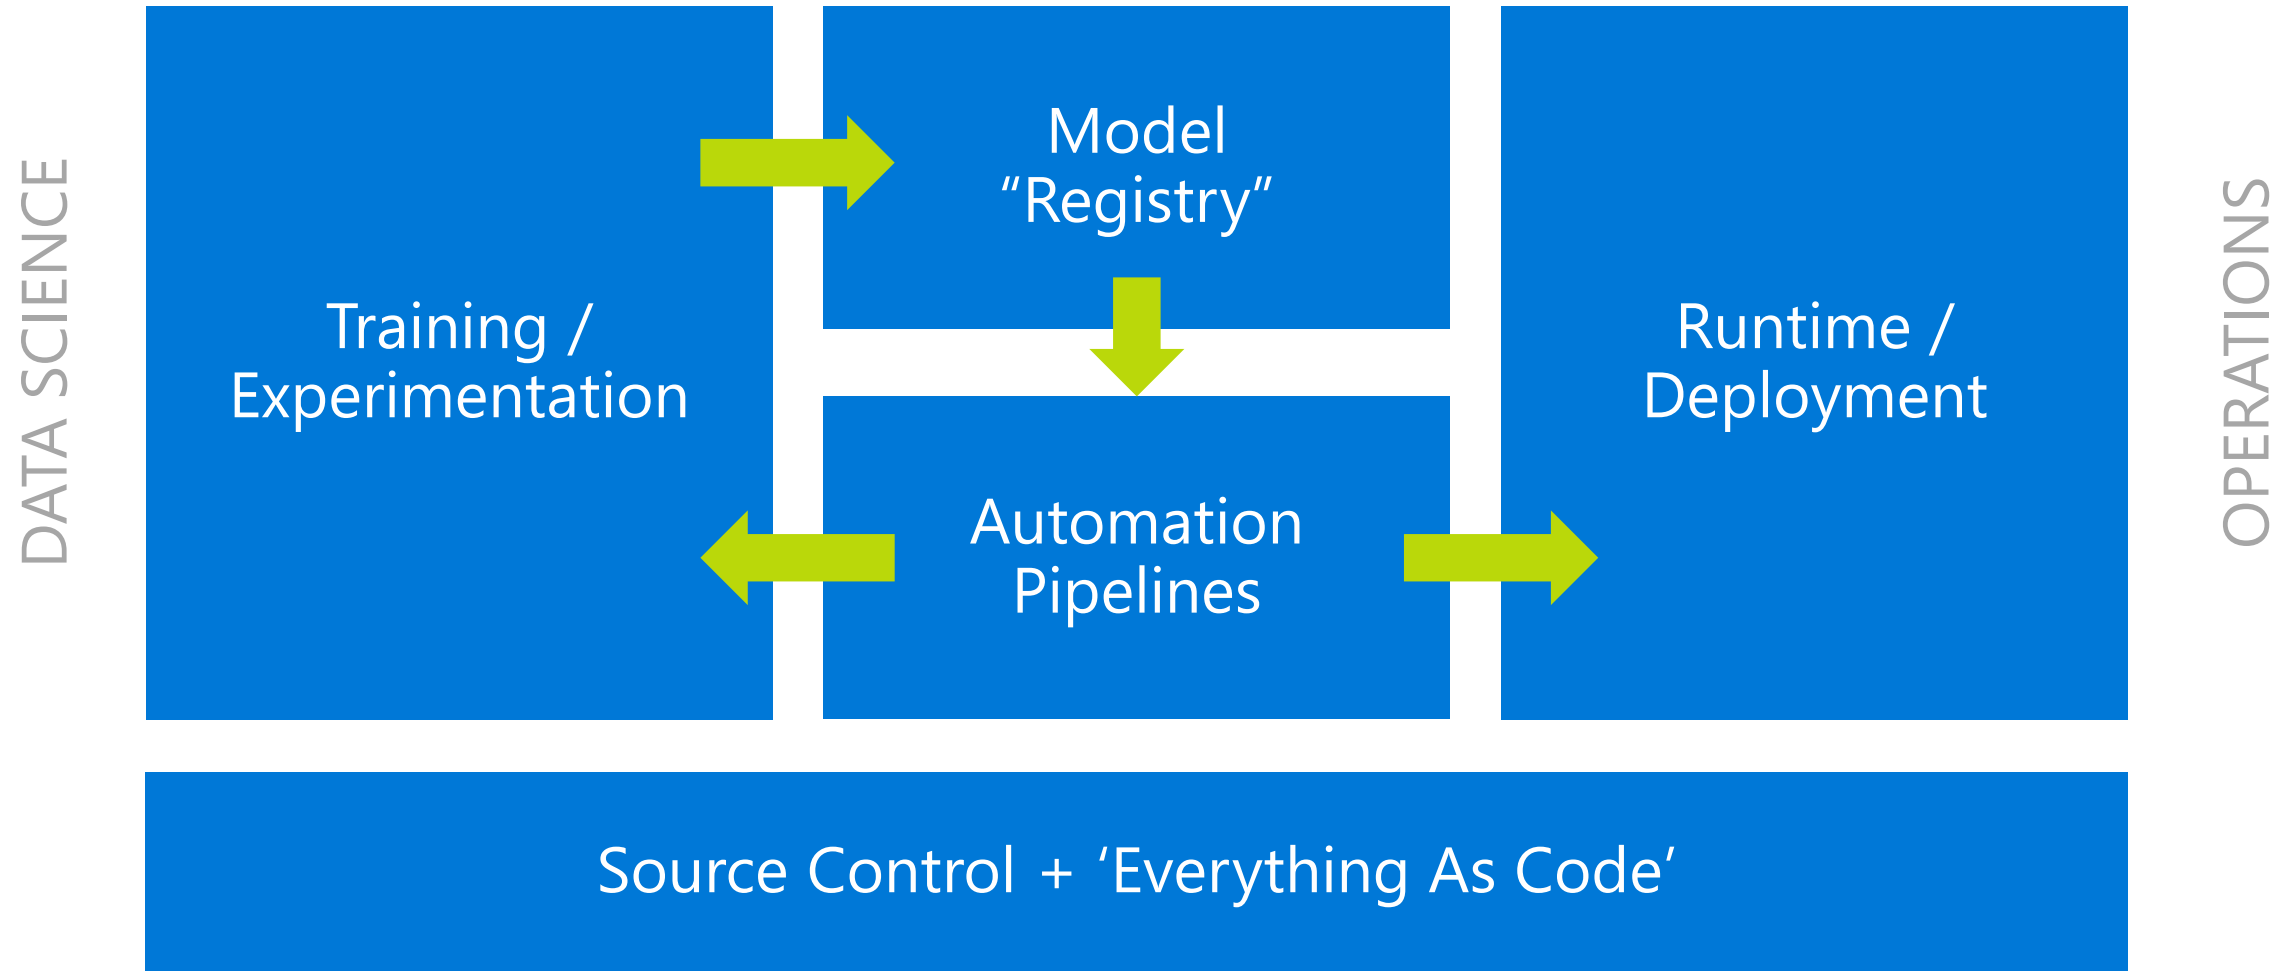
- Police recorded crime and outcomes data
- Source data as CSV - <https://data.police.uk/data>
- Build model of a given crime and region to predict – “Would you get caught?”

# Core Principals & Benefits

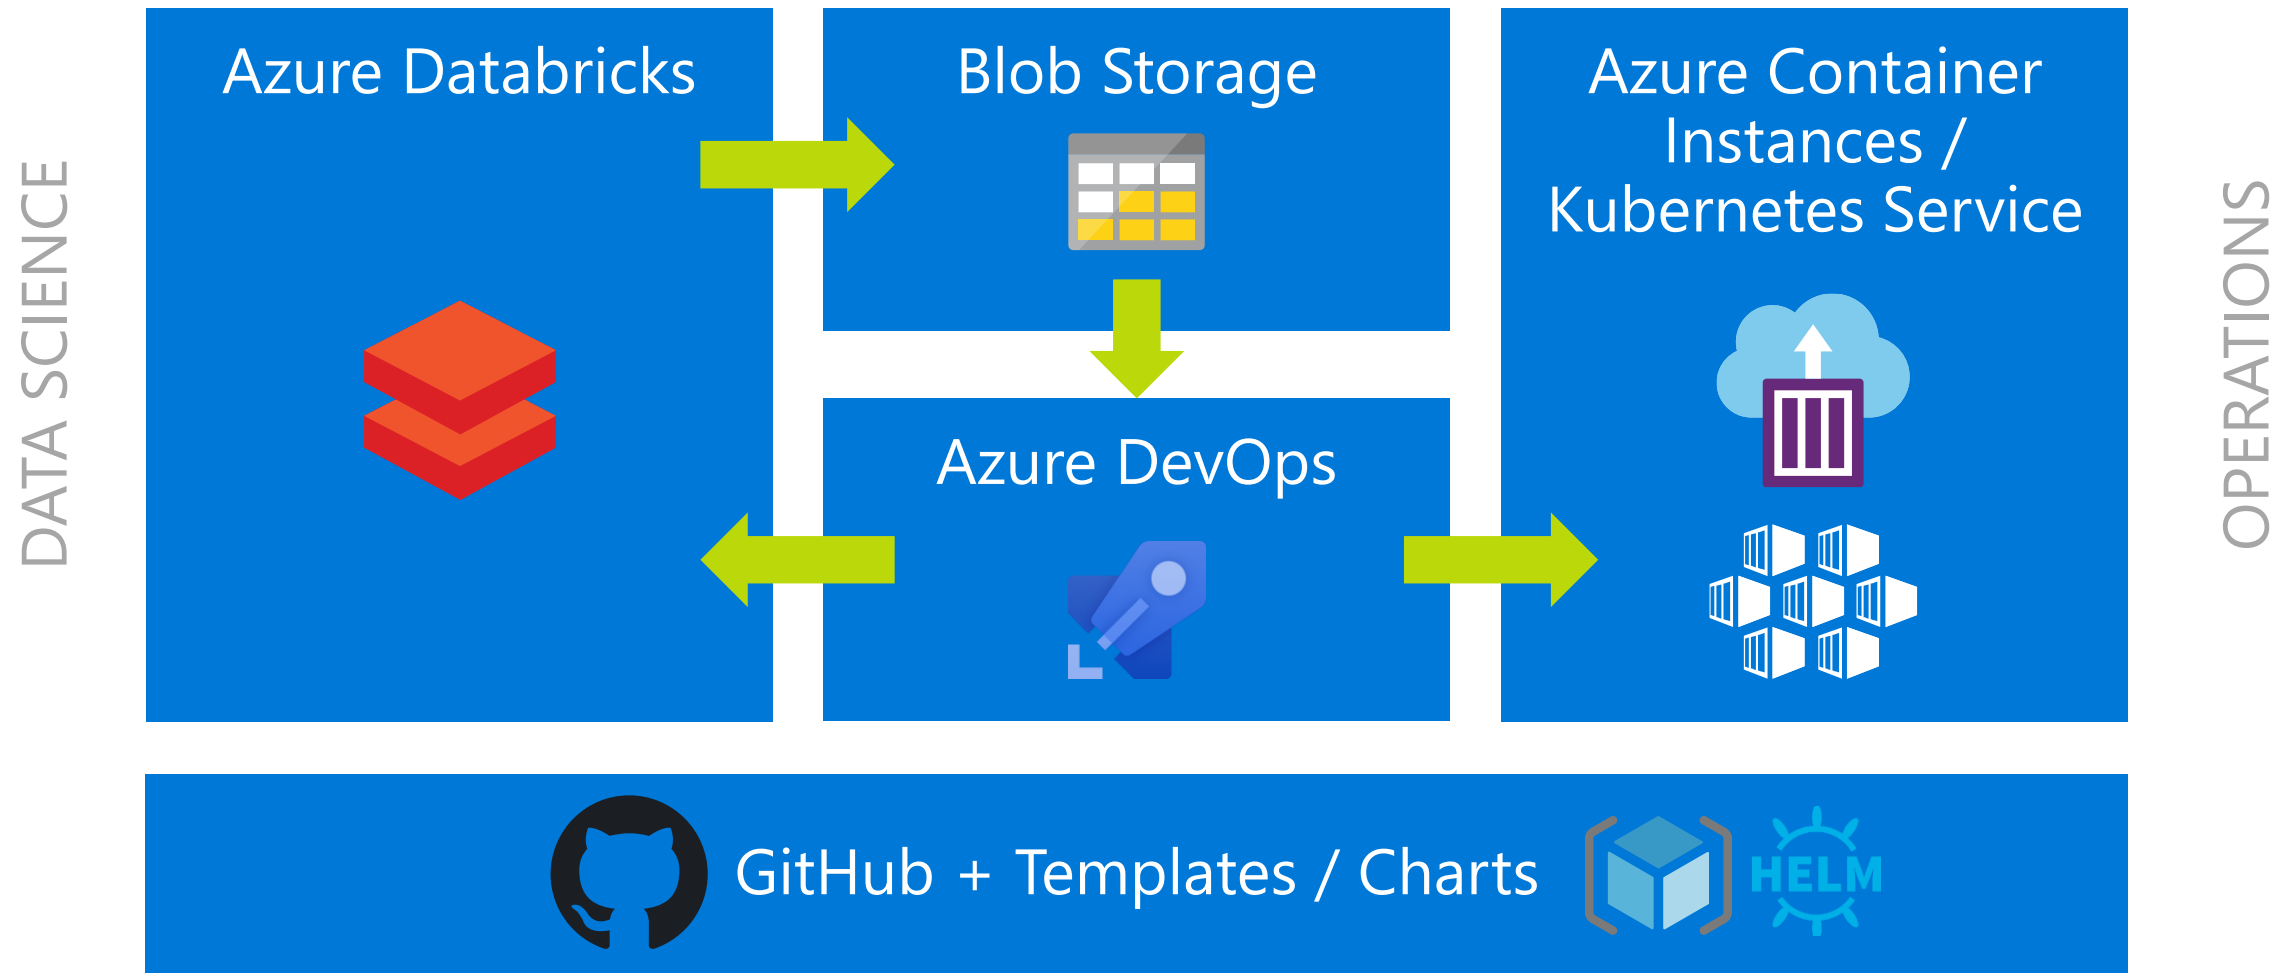
- Decouple model training experiments from operations/runtime
- Automated training, API builds & deployment
- Versioned models and API
- Config & infra as code
- Traceability



# Conceptual Building Blocks



# Conceptual Building Blocks – Project Batcomputer



# Low Level Technology Stack

Swagger



← API niceness

Gunicorn



← HTTP Server

Flask



← Web framework

Pickle



← Serialisation

Scikit-Learn



← Main ML framework

Python



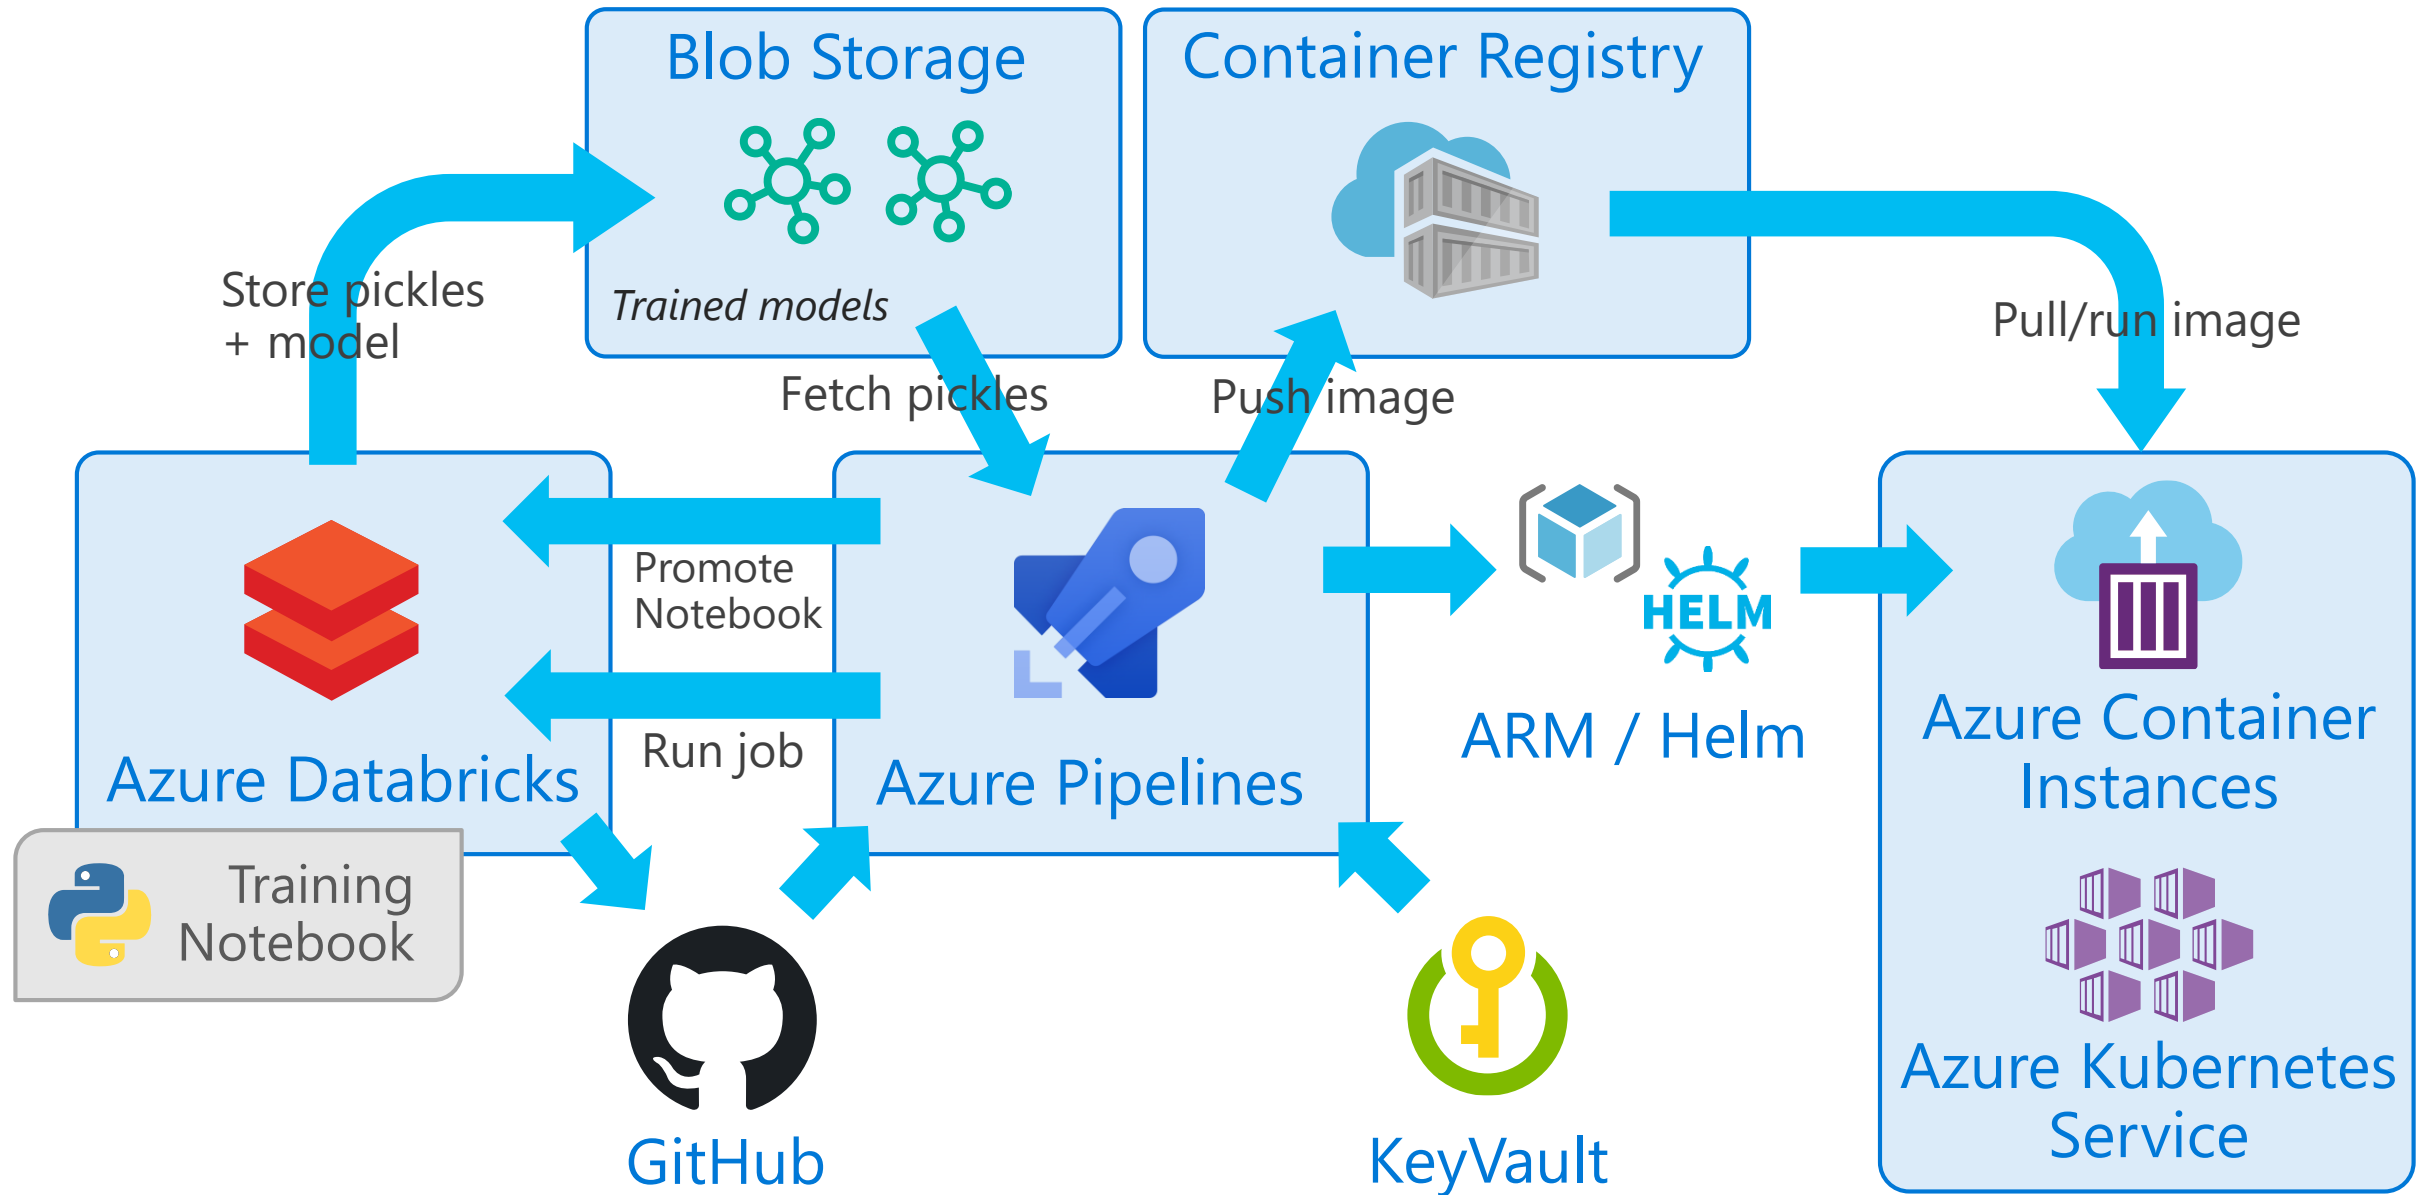
← Core language

Docker



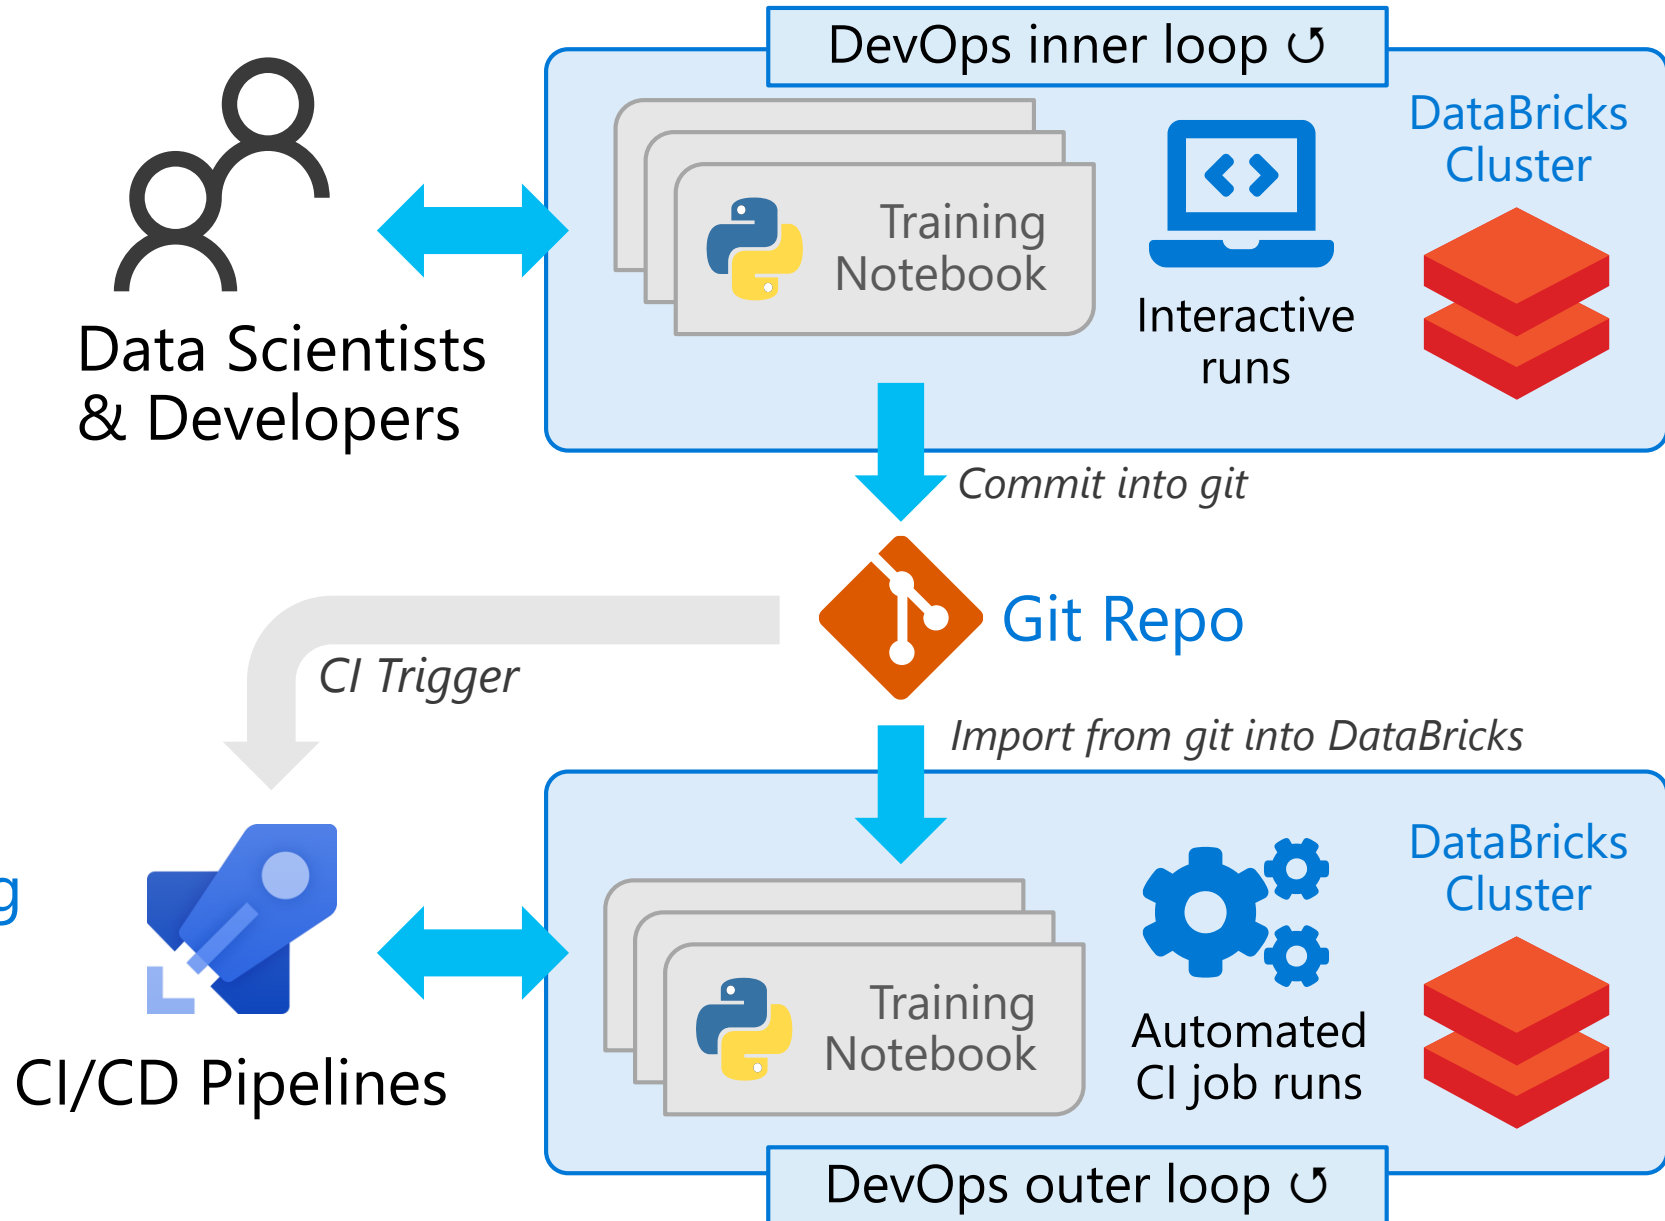
← Container Runtime

# Training & Deployment – End To End Flow




# Core DevOps Practice - Continuous Integration

- Development & experimentation
- Central shared git repo
- CI triggered training & testing job runs






# Pickling – Not Just The Model



## model.pkl

Scikit-learn model/classifier

Standard object rehydration,  
version sensitive



## lookup.pkl

Python dictionary of  
dictionaries

Mapping  
parameters/strings to num  
for predict function



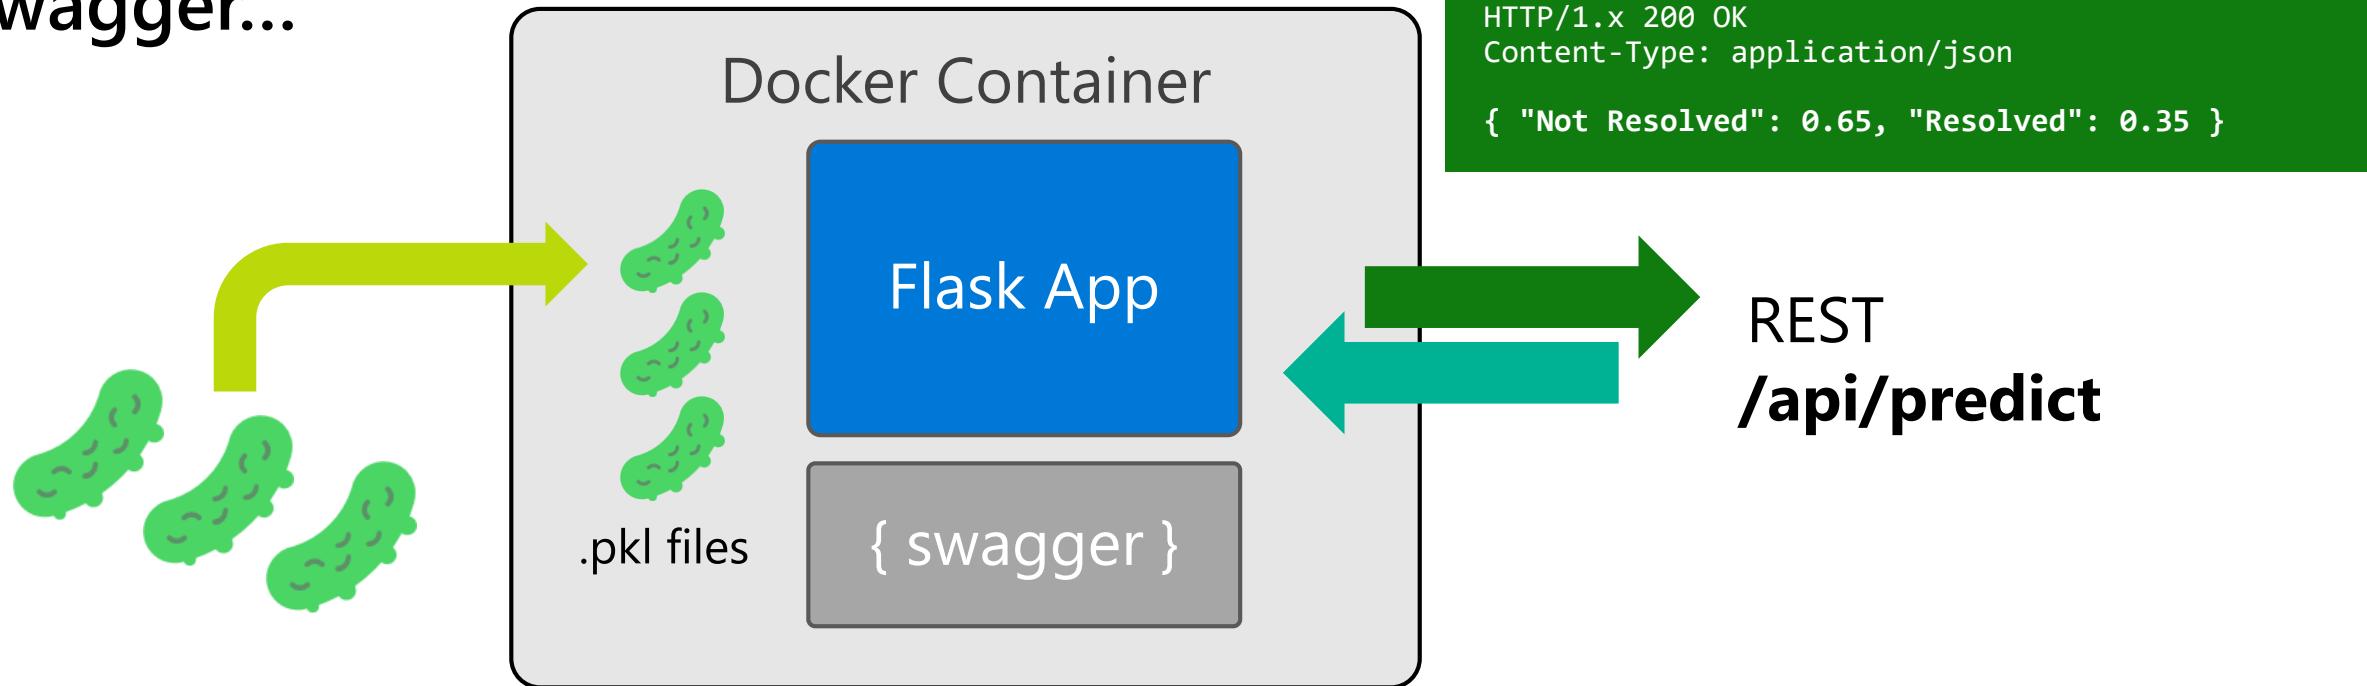
## flags.pkl

Python array

Maps output of prediction  
function to human readable  
strings or labels

# Wrapper App – Flask API

- Uses Flask web framework + Gunicorn
- Creates RESTful API for model parameters
- Consumes .pkl files
- Swagger...



# Some Decision Points


- Include model in container image or fetch at runtime?
- Wrapper app – make generic or tied into model?
- What are my API parameters?
- Convention based model registry – good enough?
- Use more robust web server than Flask?
- Automate training?



# Swagger

- We want to be RESTful
- Dynamic
  - Generated from lookup & flags pickles at runtime
- Swagger UI
  - For testing & eye candy



 **swagger**

/swagger.json

Explore

## Batcomputer API <sup>1.0.0</sup>

[ Base URL: /api ]  
</swagger.json>

REST API getting predictions from the Batcomputer ML model. Model version: 1.0.0

Schemes

HTTP

### Predictions

POST /predict

Get a prediction from the model

Parameters Try it out

Name	Description
<b>body</b> * required (body)	Request object

Example Value | Model

```
{  "offence_description": "Assault with injury",  "offence_group": "Theft offences",  "force_name": "Greater Manchester",  "offence_subgroup": "Theft from a vehicle"}
```

Parameter content type

application/json

# Semantic Versioning

## MAJOR Version

- Incompatible API or other breaking changes
- **Scoring inputs / feature changes i.e. API change**

## MINOR version

- Add functionality in a backwards-compatible manner
- **Trained using different parameters/classifiers, but API same**

## PATCH version

- Backwards-compatible bug fixes
- **Trained the model on different data**

# Versioning – Touches Everything



Also...

- Resource names in Azure controlled via ARM templates
- ACI DNS names & prefixes,  
e.g. batcomputer-2-0-8.westeurope.azurecontainer.io
- Object names in Kubernetes (pods, services), controlled via Helm chart

# Some Learnings

- Pickled Scikit-learn models are version sensitive
- Keep Python version in sync with DataBricks
- Installing numpy, scipy and scikit-learn is SLOW, pre-build base image
- Version number is the key parameter for the whole deployment process
- Writing your own wrapper isn't hard
- DataBricks has a great CLI & API
- DataBricks can be used with CI but it's not obvious



